

recommendation system architecture

<https://lucid.app/documents/embedded/9b6d3187-f2fc-490b-a958-1230637f5871>

AI Recommendation System Architecture

System Overview

This architecture diagram represents a 4-tier recommendation system for your LinkedIn-style job platform. The system is designed to efficiently match job seekers with relevant job listings based on skills, experience, and preferences.

Key Components

1. Client Application Layer

- Job Seeker Interface: Where users search for jobs and receive recommendations
- Recruiter Interface: Where employers post jobs and find candidate matches

2. API Layer (FastAPI)

- `/recommendations/jobs` : Endpoint for job recommendations to users
- `/recommendations/candidates` : Endpoint for candidate recommendations to employers
- Handles authentication, request validation, and response formatting

3. Recommendation Engine Layer

- Feature Extraction: Processes raw user and job data into comparable feature vectors
- Scoring Algorithm: Calculates match scores based on multiple factors

- Explanation Generator: Creates human-readable explanations for recommendations

4. Data Storage Layer

- User Profiles: Database storing user skills, experience, and preferences
- Job Listings: Database of current job postings and requirements
- Skill Taxonomy: Reference database for standardizing and relating skills
- Redis Cache: In-memory storage for quick access to computed recommendations

Data Flow

1. Recommendation Request Flow:

- User requests recommendations through the client interface
- Request is sent to the API layer
- API calls the recommendation engine
- Engine first checks Redis cache for existing results
- If not cached, engine:
 - Retrieves relevant data from databases
 - Extracts features from profile and job data
 - Calculates scores and generates explanations
 - Stores results in Redis cache
 - Returns formatted recommendations to client

2. Data Update Flow:

- When user profiles or job listings are updated
- Affected cache entries are invalidated
- Recommendations are recalculated on next request